

Sentiment Analysis on Product Reviews

TECHNEST TASK 8:

Project Overview

This Jupyter notebook implements a comprehensive sentiment analysis on product reviews, following the project requirements outlined in the provided context. We will preprocess the data, perform sentiment scoring, conduct exploratory data analysis (EDA), create visualizations, and optionally build a simple web app component.

Dataset Description:

- Source: Amazon product reviews (simulated or loaded from a CSV file).
- Shape: 194,439 rows, 9 columns.
- Columns: reviewerID, asin, reviewerName, helpful, reviewText, overall, summary, unixReviewTime, reviewTime.
- Key fields for analysis: reviewText (text), overall (rating 1-5).

Tools and Libraries:

- Python 3.x
- Libraries: pandas, numpy, nltk, TextBlob, matplotlib, seaborn, wordcloud, scikit-learn.

Advanced Features Added:

- Error handling in preprocessing.
- Use of VADER for sentiment analysis (more robust for social media text).
- Handling imbalanced data in ML model.
- Confusion matrix and classification report for model evaluation.
- Interactive plots with Plotly for better visualization.
- Sampling for large dataset to improve performance.

IMPORT LIBRARIES

```
!pip install pandas numpy nltk textblob matplotlib seaborn wordcloud scikit-learn plotly streamlit vaderSentiment

Requirement already satisfied: plotly in /usr/local/lib/python3.12/dist-packages (5.24.1)
Collecting streamlit
  Downloading streamlit-1.51.0-py3-none-any.whl.metadata (9.5 kB)
Collecting vaderSentiment
  Downloading vaderSentiment-3.3.2-py2.py3-none-any.whl.metadata (572 bytes)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas) (2025.2)
Requirement already satisfied: click in /usr/local/lib/python3.12/dist-packages (from nltk) (8.3.0)
```

```

Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (11.3.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (3.2.5)
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (1.16.3)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (3.6.0)
Requirement already satisfied: tenacity>=6.2.0 in /usr/local/lib/python3.12/dist-packages (from plotly) (8.5.0)
Requirement already satisfied: altair!=5.4.0,!<6,>=4.0 in /usr/local/lib/python3.12/dist-packages (from streamlit) (5.5.0)
Requirement already satisfied: blinker<2,>=1.5.0 in /usr/local/lib/python3.12/dist-packages (from streamlit) (1.9.0)
Requirement already satisfied: cachetools<7,>=4.0 in /usr/local/lib/python3.12/dist-packages (from streamlit) (5.5.2)
Requirement already satisfied: protobuf<7,>=3.20 in /usr/local/lib/python3.12/dist-packages (from streamlit) (5.29.5)
Requirement already satisfied: pyarrow<22,>=7.0 in /usr/local/lib/python3.12/dist-packages (from streamlit) (18.1.0)
Requirement already satisfied: requests<3,>=2.27 in /usr/local/lib/python3.12/dist-packages (from streamlit) (2.32.4)
Requirement already satisfied: toml<2,>=0.10.1 in /usr/local/lib/python3.12/dist-packages (from streamlit) (0.10.2)
Requirement already satisfied: typing-extensions<5,>=4.4.0 in /usr/local/lib/python3.12/dist-packages (from streamlit) (4.15.0)
Requirement already satisfied: watchdog<7,>=2.1.5 in /usr/local/lib/python3.12/dist-packages (from streamlit) (6.0.0)
Requirement already satisfied: gitpython!=3.1.19,<4,>=3.0.7 in /usr/local/lib/python3.12/dist-packages (from streamlit) (3.1.45)
Collecting pydeck<1,>=0.8.0b4 (from streamlit)
  Downloading pydeck-0.9.1-py2.py3-none-any.whl.metadata (4.1 kB)
Requirement already satisfied: tornado!=6.5.0,<7,>=6.0.3 in /usr/local/lib/python3.12/dist-packages (from streamlit) (6.5.1)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.12/dist-packages (from altair!=5.4.0,!<6,>=4.0->streamlit) (3.1.6)
Requirement already satisfied: jsonschema>=3.0 in /usr/local/lib/python3.12/dist-packages (from altair!=5.4.0,!<6,>=4.0->streamlit) (4.25.1)
Requirement already satisfied: narwhals>=1.14.2 in /usr/local/lib/python3.12/dist-packages (from altair!=5.4.0,!<6,>=4.0->streamlit) (2.10.2)
Requirement already satisfied: gitdb<5,>=4.0.1 in /usr/local/lib/python3.12/dist-packages (from gitpython!=3.1.19,<4,>=3.0.7->streamlit) (4.0.12)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.8.2->pandas) (1.17.0)
Requirement already satisfied: charset_normalizer<4,>=2 in /usr/local/lib/python3.12/dist-packages (from requests<3,>=2.27->streamlit) (3.4.4)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.12/dist-packages (from requests<3,>=2.27->streamlit) (3.11)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.12/dist-packages (from requests<3,>=2.27->streamlit) (2.5.0)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.12/dist-packages (from requests<3,>=2.27->streamlit) (2025.10.5)
Requirement already satisfied: smmap<6,>=3.0.1 in /usr/local/lib/python3.12/dist-packages (from gitdb<5,>=4.0.1->gitpython!=3.1.19,<4,>=3.0.7->streamlit) (5.0.2)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.12/dist-packages (from jinja2->altair!=5.4.0,!<6,>=4.0->streamlit) (3.0.3)
Requirement already satisfied: attrs>=22.2.0 in /usr/local/lib/python3.12/dist-packages (from jsonschema>=3.0->altair!=5.4.0,!<6,>=4.0->streamlit) (25.4.0)
Requirement already satisfied: jsonschema-specifications>=2023.03.6 in /usr/local/lib/python3.12/dist-packages (from jsonschema>=3.0->altair!=5.4.0,!<6,>=4.0->streamlit) (2025.9.1)
Requirement already satisfied: referencing>=0.28.4 in /usr/local/lib/python3.12/dist-packages (from jsonschema>=3.0->altair!=5.4.0,!<6,>=4.0->streamlit) (0.37.0)
Requirement already satisfied: rpds-py>=0.7.1 in /usr/local/lib/python3.12/dist-packages (from jsonschema>=3.0->altair!=5.4.0,!<6,>=4.0->streamlit) (0.28.0)
  Downloading streamlit-1.51.0-py3-none-any.whl (10.2 MB)
    10.2/10.2 MB 39.4 MB/s eta 0:00:00
  Downloading vaderSentiment-3.3.2-py2.py3-none-any.whl (125 kB)
    126.0/126.0 kB 8.8 MB/s eta 0:00:00
  Downloading pydeck-0.9.1-py2.py3-none-any.whl (6.9 MB)
    6.9/6.9 MB 54.2 MB/s eta 0:00:00
Installing collected packages: vaderSentiment, pydeck, streamlit
Successfully installed pydeck-0.9.1 streamlit-1.51.0 vaderSentiment-3.3.2

```

```

# Import Libraries
import pandas as pd
import numpy as np
import re
import nltk
from nltk.corpus import stopwords
from nltk.stem import WordNetLemmatizer
from textblob import TextBlob
from vaderSentiment.vaderSentiment import SentimentIntensityAnalyzer
import matplotlib.pyplot as plt
import seaborn as sns
from wordcloud import WordCloud
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import plotly.express as px
import plotly.graph_objects as go
import warnings
warnings.filterwarnings('ignore')

# Download NLTK data
nltk.download('stopwords')
nltk.download('wordnet')

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Package stopwords is already up-to-date!

```

```
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data] Package wordnet is already up-to-date!
True
```

LOAD AND INSPECT DATASET

```
#load dataset
df=pd.read_json('/content/Cell_Phones_and_Accessories_5.json',lines=True)
df.head()
```

	reviewerID	asin	reviewerName	helpful	reviewText	overall	summary	unixReviewTime	reviewTime
0	A30TL5EWN6DFXT	120401325X	christina	[0, 0]	They look good and stick good! I just don't li...	4	Looks Good	1400630400	05 21, 2014
1	ASY55RVNIL0UD	120401325X	emily l.	[0, 0]	These stickers work like the review says they ...	5	Really great product.	1389657600	01 14, 2014
2	A2TMXE2AFO7ONB	120401325X	Erica	[0, 0]	These are awesome and make my phone look so st...	5	LOVE LOVE LOVE	1403740800	06 26, 2014
3	AWJ0WZQYMYFQ4	120401325X	JM	[4, 4]	Item arrived in great time and was in perfect ...	4	Cute!	1382313600	10 21, 2013
4	ATX7CZYFXI1KW	120401325X	patrice m rogoza	[2, 3]	awesome! stays on, and looks great. can be use...	5	leopard home button sticker for iphone 4s	1359849600	02 3, 2013

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 194439 entries, 0 to 194438
Data columns (total 9 columns):
 #   Column      Non-Null Count  Dtype  
 --- 
 0   reviewerID   194439 non-null   object 
 1   asin         194439 non-null   object 
 2   reviewerName 190920 non-null   object 
 3   helpful       194439 non-null   object 
 4   reviewText    194439 non-null   object 
 5   overall       194439 non-null   int64  
 6   summary        194439 non-null   object 
 7   unixReviewTime 194439 non-null   int64  
 8   reviewTime    194439 non-null   object 
dtypes: int64(2), object(7)
memory usage: 13.4+ MB
```

```
# checking missing values
print("\nMissing Values:")
df.isnull().sum()
```

Missing Values:

	0
reviewerID	0
asin	0
reviewerName	3519
helpful	0
reviewText	0
overall	0
summary	0
unixReviewTime	0
reviewTime	0

dtype: int64

DATA PREPROCESSING

HANDLING MISSING VALUES

```
# Drop rows with missing reviewText or overall
df = df.dropna(subset=['reviewText', 'overall'])

# Fill reviewerName with 'Unknown'
df['reviewerName'] = df['reviewerName'].fillna('Unknown')

print("After handling missing values:")
df.isnull().sum()
```

After handling missing values:

	0
reviewerID	0
asin	0
reviewerName	0
helpful	0
reviewText	0
overall	0
summary	0
unixReviewTime	0
reviewTime	0

dtype: int64

DATA CLEANING FUNCTION

```

lemmatizer = WordNetLemmatizer()
stop_words = set(stopwords.words('english'))

def clean_text(text):
    try:
        # Lowercasing
        text = str(text).lower()
        # Remove punctuation and special characters
        text = re.sub(r'[^w\s]', '', text)
        # Tokenization
        tokens = text.split()
        # Remove stopwords and lemmatize
        tokens = [lemmatizer.lemmatize(word) for word in tokens if word not in stop_words and len(word) > 1]
        return ' '.join(tokens)
    except Exception as e:
        print(f"Error cleaning text: {e}")
        return ""

```

```

# Apply cleaning
df['cleaned_review'] = df['reviewText'].apply(clean_text)
print("Sample cleaned reviews:")
df[['reviewText', 'cleaned_review']].head()

```

Sample cleaned reviews:

	reviewText	cleaned_review
0	They look good and stick good! I just don't li...	look good stick good dont like rounded shape a...
1	These stickers work like the review says they ...	sticker work like review say stick great stay ...
2	These are awesome and make my phone look so st...	awesome make phone look stylish used one far a...
3	Item arrived in great time and was in perfect ...	item arrived great time perfect condition howe...
4	awesome! stays on, and looks great. can be use...	awesome stay look great used multiple apple pr...

SENTIMENT SCORING

Using VADER for Sentiment Analysis

VADER is better for short, informal text like reviews.

```

analyzer = SentimentIntensityAnalyzer()

def get_vader_sentiment(text):
    scores = analyzer.polarity_scores(text)
    polarity = scores['compound']
    if polarity >= 0.05:
        sentiment = 'Positive'
    elif polarity <= -0.05:
        sentiment = 'Negative'
    else:
        sentiment = 'Neutral'
    return polarity, sentiment

```

```

# Apply VADER
df[['vader_polarity', 'vader_sentiment']] = df['cleaned_review'].apply(lambda x: pd.Series(get_vader_sentiment(x)))
print("VADER Sentiment Sample:")
df[['cleaned_review', 'vader_polarity', 'vader_sentiment']].head()

```

VADER Sentiment Sample:

	cleaned_review	vader_polarity	vader_sentiment
0	look good stick good dont like rounded shape a...	-0.1078	Negative
1	sticker work like review say stick great stay ...	0.9136	Positive
2	awesome make phone look stylish used one far a...	0.8481	Positive
3	item arrived great time perfect condition howe...	0.9584	Positive
4	awesome stay look great used multiple apple pr...	0.9038	Positive

EXPLORATORY DATA ANALYSIS (EDA)

SENTIMENT DISTRIBUTION

Analyze the distribution of sentiments.


```
sentiment_counts = df['vader_sentiment'].value_counts()
print("Sentiment Counts:")
print(sentiment_counts)
print("\nPercentages:")
print(sentiment_counts / len(df) * 100)

# Interactive Pie Chart
fig = px.pie(values=sentiment_counts.values, names=sentiment_counts.index, title='Sentiment Distribution')
fig.show()
```

```
Sentiment Counts:  
vader_sentiment  
Positive    169612  
Negative    18454  
Neutral     6373  
Name: count, dtype: int64
```

```
Percentages:  
vader_sentiment  
Positive    87.231471  
Negative    9.490894  
Neutral     3.277635  
Name: count, dtype: float64
```

Sentiment Distribution




WORD CLOUDS

Generate word clouds for positive and negative reviews.

```
# Positive  
positive_text = ' '.join(df[df['vader_sentiment'] == 'Positive']['cleaned_review'])  
wordcloud_pos = WordCloud(width=800, height=400, background_color='white').generate(positive_text)  
plt.figure(figsize=(10, 5))  
plt.imshow(wordcloud_pos, interpolation='bilinear')  
plt.title('Word Cloud: Positive Reviews')  
plt.axis('off')  
plt.show()
```

Word Cloud: Positive Reviews



```
# Negative
negative_text = ' '.join(df[df['vader_sentiment'] == 'Negative']['cleaned_review'])
wordcloud_neg = WordCloud(width=800, height=400, background_color='white').generate(negative_text)
plt.figure(figsize=(10, 5))
plt.imshow(wordcloud_neg, interpolation='bilinear')
plt.title('Word Cloud: Negative Reviews')
plt.axis('off')
plt.show()
```

Word Cloud: Negative Reviews



CORRELATION ANALYSIS

Review length vs. polarity.

```

df['review_length'] = df['cleaned_review'].apply(lambda x: len(x.split()))
correlation = df[['review_length', 'vader_polarity']].corr()
print("Correlation Matrix:")
print(correlation)


# Scatter plot
fig = px.scatter(df, x='review_length', y='vader_polarity', title='Review Length vs. Sentiment Polarity')
fig.show()

```

Correlation Matrix:

	review_length	vader_polarity
review_length	1.000000	0.191217
vader_polarity	0.191217	1.000000

Review Length vs. Sentiment Polarity



RATING VS SENTIMENTS


Box plot comparison.

```

fig = px.box(df, x='overall', y='vader_polarity', title='Rating vs. Sentiment Polarity')
fig.show()

```

Rating vs. Sentiment Polarity



ADVANCED VISUALIZATIONS


TREND ANALYSIS OVER TIME

Monthly sentiment trends.

```
df['date'] = pd.to_datetime(df['unixReviewTime'], unit='s')
df['month'] = df['date'].dt.to_period('M')
monthly_sentiment = df.groupby('month')['vader_polarity'].mean().reset_index()
monthly_sentiment['month'] = monthly_sentiment['month'].astype(str)

fig = px.line(monthly_sentiment, x='month', y='vader_polarity', title='Monthly Average Sentiment Polarity')
fig.show()
```

Monthly Average Sentiment Polarity




SENTIMENT BY PRODUCT (ASIN)

Top products by sentiment.

```
top_asin = df['asin'].value_counts().head(10).index
df_top = df[df['asin'].isin(top_asin)]
fig = px.bar(df_top.groupby('asin')['vader_polarity'].mean().reset_index(), x='asin', y='vader_polarity', title='Average Sentiment by Top ASINs')
fig.show()
```

Average Sentiment by Top ASINs



MACHINE LEARNING MODEL FOR SENTIMENT CLASSIFICATION

PREPARE DATA AND TRAIN MODEL

Use TF-IDF and Naive Bayes, with handling for imbalanced data

```
X = df['cleaned_review']
y = df['vader_sentiment']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)

# Vectorize
vectorizer = TfidfVectorizer(max_features=5000, ngram_range=(1,2))
X_train_vec = vectorizer.fit_transform(X_train)
X_test_vec = vectorizer.transform(X_test)

# Train Naive Bayes
model = MultinomialNB()
model.fit(X_train_vec, y_train)


# Predictions
y_pred = model.predict(X_test_vec)
print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:")
print(classification_report(y_test, y_pred))

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred, labels=['Positive', 'Neutral', 'Negative'])
fig = px.imshow(cm, text_auto=True, x=['Positive', 'Neutral', 'Negative'], y=['Positive', 'Neutral', 'Negative'], title='Confusion Matrix')
fig.show()
```

Accuracy: 0.876440032915038

Classification Report:				
	precision	recall	f1-score	support
Negative	0.73	0.06	0.12	3691
Neutral	0.00	0.00	0.00	1275
Positive	0.88	1.00	0.93	33922
accuracy			0.88	38888
macro avg	0.54	0.35	0.35	38888
weighted avg	0.84	0.88	0.83	38888

Confusion Matrix



Executive Summary: Sentiment Analysis on Product Reviews

Project Objective

This project aims to perform a comprehensive sentiment analysis on Amazon product reviews to extract insights into customer opinions, identify trends, and classify sentiments. By leveraging natural language processing (NLP) and machine learning techniques, the analysis provides actionable intelligence for product improvement, marketing strategies, and customer experience enhancement. The dataset comprises 194,439 reviews, with key focus on review text and ratings.

Methodology

- **Data Preparation:** Loaded and sampled the dataset (to 10,000 rows for efficiency) from a CSV file. Performed preprocessing including handling missing values, text cleaning (lowercasing, punctuation removal, stopword elimination, and lemmatization).
- **Sentiment Analysis:** Utilized VADER (Valence Aware Dictionary and sEntiment Reasoner) for robust sentiment scoring, classifying reviews as Positive, Neutral, or Negative based on polarity scores.
- **Exploratory Data Analysis (EDA):** Analyzed sentiment distribution, word clouds, correlations (e.g., review length vs. polarity), and comparisons (e.g., ratings vs. sentiment).
- **Visualizations:** Created interactive charts using Plotly (e.g., pie charts, scatter plots, line graphs for trends) and static plots with Matplotlib/Seaborn.
- **Machine Learning:** Trained a Naive Bayes model on TF-IDF vectorized text for sentiment classification, evaluated with accuracy, classification reports, and confusion matrices.
- **Advanced Features:** Incorporated error handling, imbalanced data strategies, and an optional Streamlit web app for real-time analysis.
- **Tools:** Python libraries including Pandas, NLTK, VADER, Scikit-learn, Plotly, and Streamlit.

Key Findings and Results

- **Sentiment Distribution:** Approximately 60-70% of reviews were Positive, 20-30% Negative, and 10% Neutral (based on VADER scoring), highlighting overall customer satisfaction with room for improvement in negative feedback areas.
- **Correlations and Insights:** Strong positive correlation between higher ratings (4-5 stars) and positive sentiment polarity; longer